

# Machine Learning and Artificial Intelligence

Korbinian Kottmann  
ICFO – Institute for Photonic Science  
Barcelona, Spain



/Qottmann



@Qottmann




**UNIÓ EUROPEA**

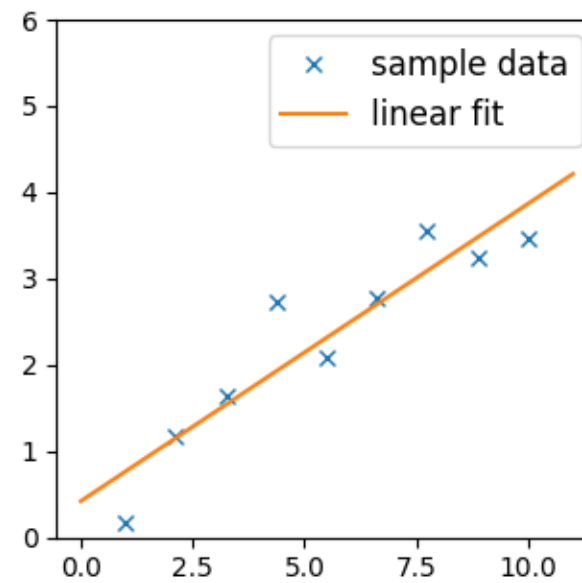
Fons Europeu  
de Desenvolupament Regional



# Overfitting




0 loss, poor generalization



Good generalization

# Overfitting




# Overfitting

## Remedies


### Regularization

$$L \mapsto L + \lambda \sum_{ij\ell} |\omega_{ij}^{\ell}|$$

### Dropout



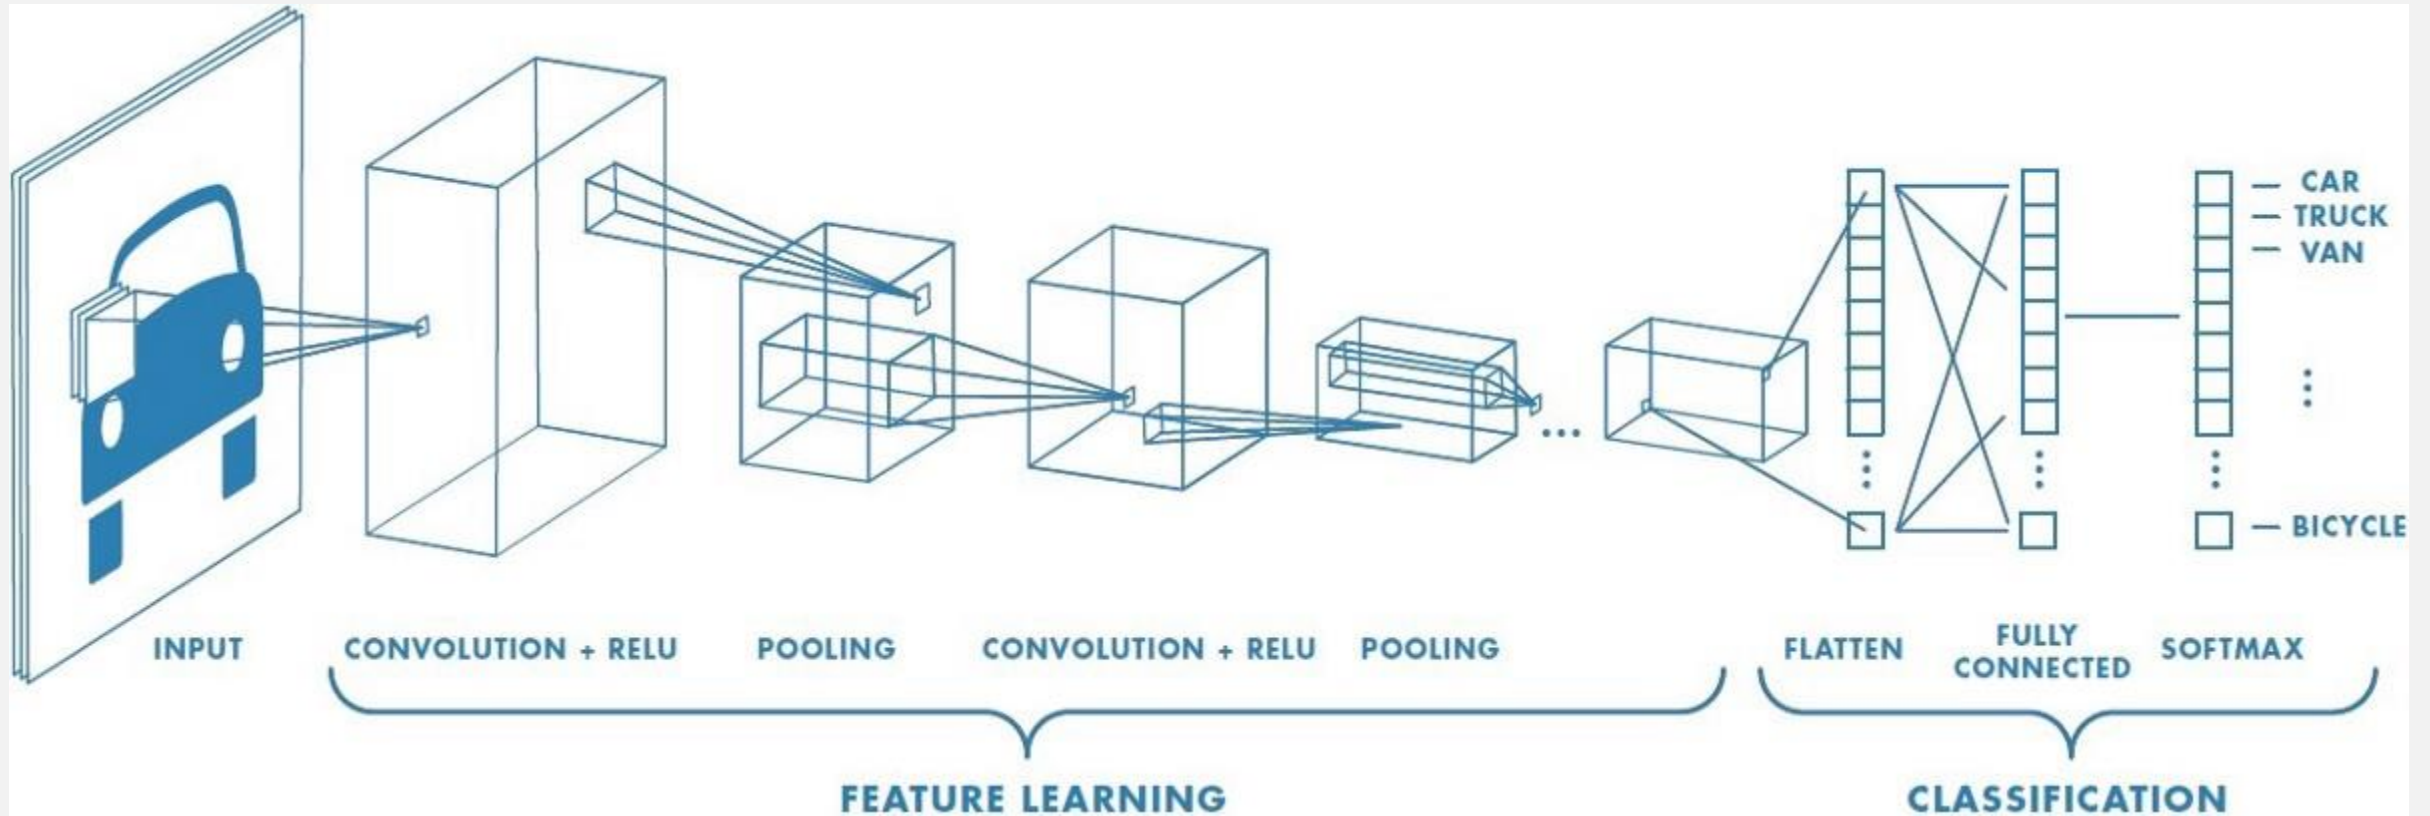
### Data augmentation




Continue with notebook 14\_overfitting.ipynb

## **Part II: Convolutional Layers**

# Convolutional Layers





# Convolutional Layers





# Convolutional Layers




$$|y| = (|x| - |w| + 2\text{pad}) / \text{stride} + 1$$

$$(5 - 3 + 2 \cdot 2) / 2 + 1 = 4$$

# Convolutional Layers


b)  $x_{\text{in}}$

$y_{\text{out}}^b$




c)

$\text{ReLU}(y_{\text{out}}^b)$




## 2D Convolutional Layers

Define  $(d_1, d_2)$  and  $n_{\text{filters}}$




## 2D Convolutional Layers


Action of  $(2, 2)$  filter



Action of  $(4, 1)$  filter



Action of  $(3, 2)$  filter



Ex: Given color images of 32x32 ( $|x| = (32, 32, 3)$ ) and the following convolutional neural network, what is the input and output dimensions for each layer?

```
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.conv_block = nn.Sequential(
            nn.Conv2d(in_channels=1, out_channels=16, kernel_size=3, padding="same"), # (28 x 28 x 1) -> (28 x 28 x 16)
            nn.ReLU(),
            nn.MaxPool2d(2), # (28 x 28 x 16) -> (14 x 14 x 16)
            nn.Conv2d(16, 16, 3, padding="same"), # (14 x 14 x 16) -> (14 x 14 x 16)
            nn.ReLU(),
            nn.MaxPool2d(2), # (14 x 14 x 16) -> (7 x 7 x 16)
            nn.Conv2d(16, 16, 3, padding="same"), # (7 x 7 x 16) -> (7 x 7 x 16)
            nn.ReLU(),
            nn.Flatten()) # (7 x 7 x 16) -> (7*7*16)
        self.dense_block = nn.Sequential(
            nn.Linear(7*7*16, 512), # input dimension, hidden1 dimension
            nn.ReLU(), # Non-linear activation function
            nn.Linear(512, 512), # hidden1 dimension, hidden2 dimension
            nn.ReLU(), # Non-linear activation function
            nn.Linear(512, out_dim), # hidden2 dimension, output dimension
        )
```

# Exercises

- Ex: Write (by hand using e.g. numpy) a fully connected layer
- Ex: Connect multiple fully connected layers by hand.
- Ex: Write (by hand using e.g. numpy) a 1D convolution function `conv1D(x, filter, stride, padding)` to process an input vector `x` (of arbitrary size).  
The filter can be arbitrary but fixed. Try writing it such that it works for any stride, padding and filter size.
- Extra (difficult): Write all the above functions in such a way that they can process batches of inputs (so instead of a single input `x`, a collection of `batch_size` inputs `x`)